

STK 880: Bayesian Statistical Modelling and Computing using R, JAGS and (Stan)

Working with JAGS and R

Dr. Rianne Jacobs

University of Groningen, The Netherlands

14-18 September 2018

Run JAGS from R

- Install and load packages

```
> install.packages("runjags")  
> install.packages("coda")  
> library("runjags")  
> library("coda")
```

- Set working directory

```
> setwd("C:/Documents/STK880_2018/Jags_Tutorial/")
```

- Load the data and check

```
> ALP <- read.table("data/ALP.txt", header = TRUE)  
> head(ALP)  
> str(ALP)  
> summary(ALP)
```

Run JAGS from R

- Data transformations

```
> ALP$alkfos_tr <- 100/sqrt(ALP$alkfos)
> hist(ALP$alkfos_tr)
```

- Create model file and check

```
> file.show("osteo.model.txt")
```

- Prepare data and gather in a list

```
> alp <- ALP[ALP$artikel==0,"alkfos_tr"]
> n.obs <- length(alp)
> data.list <- list(y = alp, n.obs = n.obs)
```

Run JAGS from R

- Prepare initial values

- Explicitly for each chain

```
> inits.list <- list(  
+ list(mu = 5.5, tau = 2.5),  
+ list(mu = 5.2, tau = 2.2))
```

- As a function

```
> inits.list <- function()(list(mu = rnorm(1,5.5,1),  
+ tau = runif(1,2,3)))
```

- Parameters to monitor - those for which you want a posterior sample

```
> parameters <- c("mu","sig2")
```

Run JAGS from R

- Run the `run.jags` function in R

```
> post.sim <- run.jags(model = "NormalDistribution",  
+ data = data.list, inits = inits.list,  
+ monitor = parameters,  
+ n.chains = 2, burnin = 500, thin = 1, sample = 5000)
```

- Provide the model, data, initial values and parameter vector to monitor
- number of chains (`n.chains = 2`), number of burnin iterations (`burnin = 500`), thinning factor (`thin = 1`), number of iterations per chain (`sample = 5000`)
- Produce plots

```
> plot(post.sim, plot.type = c("trace"), vars = "mu")
```
- Produce summary statistics

```
> summary(post.sim)
```

Run JAGS from R

- Convert `run.jags` object to `mcmc` object

```
> samples <- as.mcmc(post.sim)
```

- CODA plots and functions

```
> par(mfrow=c(2,2)) # plot figures in 2x2 format
```

```
> traceplot(samples) # trace plots
```

```
> cumuplot(samples,ask=FALSE) # running mean plots
```

```
> acfplot(samples) # autocorrelation function plot
```

```
> autocorr(samples) # autocorrelation values
```

```
> crosscorr.plot(samples) # cross-correlation output
```

```
> densplot(samples) # density plots of the marginal s
```

```
> effectiveSize(samples) # effective size
```

```
> HPDinterval(samples) # HPD intervals
```

Run JAGS from R

- CODA convergence diagnostics
 - > `gelman.diag(osteo.mcmc)`
 - > `gelman.plot(osteo.mcmc,ask=FALSE)`
 - > `geweke.diag(osteo.mcmc)`
 - > `geweke.plot(osteo.mcmc,ask=FALSE)`
- Interpret results

Adaptation and burnin

- When a model is initialized, it may be in adaptive mode
- Samplers used by the model may modify their behaviour for increased efficiency
- Sequence generated by an adapting sampler is no longer a Markov chain - not guaranteed to converge to the target distribution
- Sufficient number of iterations must take place after the adaptive phase to ensure successful burnin
- `run.jags` default: `adapt = 1000`, `burnin = 4000`

Choice of priors

- Incorporating prior knowledge
- Unique feature for Bayesian approach
- But might introduce subjectivity - Good to include sensitivity analysis
- Different kinds of priors
 - Conjugate
 - Noninformative
 - Informative

Conjugate priors

Table 5.2 Common members of the exponential family and their associated (natural) conjugate prior.

Exponential family member		Parameter	Conjugate prior
Univariate case			
Discrete distributions			
Bernoulli	$\text{Bern}(\theta)$	θ	$\text{Beta}(\alpha_0, \beta_0)$
Binomial	$\text{Bin}(n, \theta)$	θ	$\text{Beta}(\alpha_0, \beta_0)$
Negative binomial	$\text{NB}(k, \theta)$	θ	$\text{Beta}(\alpha_0, \beta_0)$
Poisson	$\text{Poisson}(\theta)$	θ	$\text{Gamma}(\alpha_0, \beta_0)$
Continuous distributions			
Normal-variance fixed	$N(\mu, \sigma^2) - \sigma^2$ fixed	μ	$N(\mu_0, \sigma_0^2)$
Normal-mean fixed	$N(\mu, \sigma^2) - \mu$ fixed	μ	$\text{IG}(\alpha_0, \beta_0)$
Normal*	$N(\mu, \sigma^2)$	μ, σ^2	$\text{Inv-}\chi^2(\nu_0, \tau_0^2)$
			$\text{NIG}(\mu_0, \kappa_0, a_0, b_0)$
Exponential	$\text{Exp}(\lambda)$	λ	$\text{N-Inv-}\chi^2(\mu_0, \kappa_0, \nu_0, \tau_0^2)$
			$\text{Gamma}(\alpha_0, \beta_0)$

Conjugate priors

Multivariate case			
Discrete distributions			
Multinomial	Mult(n, θ)	θ	Dirichlet(α_0)
Continuous distributions			
Normal-covariance fixed	N(μ, Σ)- Σ fixed	μ	N(μ_0, Σ_0)
Normal-mean fixed	N(μ, Σ)- μ fixed	Σ	IW(Λ_0, ν_0)
Normal*	N(μ, Σ)	μ, Σ	NIW($\mu_0, \kappa_0, \nu_0, \Lambda_0$)

- Conjugacy: Posterior same type as prior
- Computational advantage in JAGS
- Interpretation: likelihood of historical data can be easily turned into a conjugate prior

Noninformative priors

- Expresses no knowledge
- non-subjective, objective, default, reference, weak, diffuse, flat, conventional, minimally informative, ...
- Noninformative on what scale?
- Always good to do a sensitivity analysis
- Ignorance cannot be expressed mathematically
- Jeffrey's priors most popular (pgs 115-117 (Lesaffre, Lawson, 2012))
- In practice
 - Locally uniform, vague or weak
 - e.g. $\beta_1 \sim N(0, 100^2)$
 - e.g. $1/\sigma^2 \sim \text{Gamma}(0.001, 0.001)$

Informative priors

- Examples
 - Based on historical studies
 - Literature
 - Expert knowledge
- Make sure that informative prior has a sound scientific basis

The JAGS model

Series of relations inside a block delimited by curly brackets { and } and preceded by the keyword model

```
model {  
  for (i in 1:N) {  
    Y[i] ~ dnorm(mu[i], tau)  
    mu[i] <- alpha + beta * (x[i] - x.bar)  
  }  
  x.bar <- mean(x)  
  alpha ~ dnorm(0.0, 1.0E-4)  
  beta ~ dnorm(0.0, 1.0E-4)  
  sigma <- 1.0/sqrt(tau)  
  tau ~ dgamma(1.0E-3, 1.0E-3)  
}
```

The JAGS model

- Each relation defines a node in the model
- Nodes in the model form a directed acyclic graph
- Two types of relations
 - Stochastic relation (\sim) defines a stochastic node - random variable in the model
 - Deterministic relation ($< -$) defines a deterministic node - value is determined exactly by the values of its parents
- See JAGS manual for available distributions
- NOTE: $N(\mu, \tau)$, $\tau = 1/\sigma^2$ - precision
- Use of scalars, vectors, matrices, arrays as in R

The JAGS model

- For loops

```
for (i in 1:3) {  
  Y[i] ~ dnorm(mu, tau)  
}
```

- is equivalent to

```
Y[1] ~ dnorm(mu, tau)  
Y[2] ~ dnorm(mu, tau)  
Y[3] ~ dnorm(mu, tau)
```


Compiling the JAGS model

Compilation may fail for a number of reasons

- The model uses a function or distribution that has not been defined in any of the loaded modules
- The graph contains a directed cycle. These are forbidden in JAGS.
- A node is used in an expression but never defined by a relation or supplied with the data.

Other runjags functions and options

- `run.jags`
- `extend.jags` to continue running the model
- `autorun.jags` takes control of the run-length of the MCMC process and runs the chain until convergence
- parallel processing
 - `method = "parallel"`
 - runs each chain on a different core
- modules
 - `modules = glm`
 - The `glm` module implements samplers for efficient updating of generalized linear mixed models.
 - block updating - no need to centre predictor variables